

Random maze generation by recursive division, and shortest-path search to the exit

Mohamed Ali Tewfik Hamlil

M1 Artificial Intelligence, Université Grenoble Alpes

Algorithmic Problem Solving — APP1

September 12, 2026

Abstract

A rectangular area of $W \times H$ cells is laid out as a maze by recursive division: a full wall across a dimension, exactly one door in it, and the same rule inside each of the two halves. The requirements force the result to be a spanning tree of the grid graph, and the generator returns it *twice*: as a graph of cells joined by doors, and as a binary tree recording which division made which sub-maze. Each structure has its own way out. The graph is searched with Dijkstra's algorithm, every door costing one step; the tree is read with a lowest-common-ancestor descent, which never touches a cell. Both return the same route — the maze is a tree, so there is only one — so what separates them is how much they examine, and the tree wins by a factor of thirty. A C implementation measures every claim, and the closing section puts costs on all of it and against the literature.

Contents

1	Problem analysis and formulation	3
1.1	Statement	3
1.2	From a drawing of walls to a graph	3
1.3	What the specification forces E to be	4
1.4	What a route costs	5
1.5	Constraints on the exit	5
1.6	The two problems, restated on the graph	5
1.7	Conventions	6
1.8	Scope of the deliverable	6
2	Maze generation, and the two structures it returns	6
2.1	Two outputs, not one	6
2.2	The node record	7
2.3	Regions, named by two corners	7
2.4	The recursion stops at a corridor	7
2.5	One division contributes exactly one door	8
2.6	The second output: the binary tree	8
2.7	One division step	9
2.8	The algorithm	10
2.9	Why the result is a spanning tree	10
2.10	What generation returns	11
3	Solving the graph output: Dijkstra	11
3.1	Dijkstra on the maze	11
3.2	The two distances to the exit	12
3.3	Two ways to use an estimate	12
3.4	What tie-breaking can and cannot do	12
3.5	The algorithm	13
4	Solving the tree output: the common ancestor	14
4.1	Addresses	14
4.2	The lemma that does the work	14
4.3	The algorithm	15
4.4	Why not walk the tree greedily	15
5	Implementation in C	16
5.1	What was built	16
5.2	What is checked	17
6	Time analysis and comparison	17
6.1	Size of the two structures	18
6.2	Cost of generation	18
6.3	Cost of solving the graph	18
6.4	Cost of solving the tree	18
6.5	The comparison	19
6.6	Why the heuristic is weak here, and the tree is not	20
6.7	Where recursive division sits	20

1 Problem analysis and formulation

1.1 Statement

A rectangular area of $W \times H$ unit cells is to be laid out as a maze.

Three things are fixed. The layout is drawn **at random**, so W and H do not determine it. The area is **enclosed** and has a single **exit**. It is entered at an arbitrary cell.

Two problems follow. Both are settled on paper first, then analysed for cost.

1. Given the width and the height of the initial area, generate a maze.
2. Find the path from any position in the maze to the exit. The specification states it is unique, hence shortest.

The construction is prescribed.

Divide the allocated space with a full **wall** across its longest dimension. Leave exactly one **door** in that wall. Apply the same rule to each of the two spaces this creates.

Figure 1 is the entire rule. Everything below follows from it.




Figure 1: The construction rule. Each application turns one region into two regions of the same kind joined by a single door, which is what makes this a divide-and-conquer problem and the decomposition a binary tree.

What “at random” binds. Half the output, not all of it.

The cells are fixed by W and H . The same WH positions appear in every run, and their coordinates are public the moment the dimensions are.

What must be unpredictable is **which doors were kept** — the edge set.

So the requirement falls on the draws, not on the storage. A rule that takes the wall from the geometry, the midpoint say, makes the maze a function of W and H . Anyone holding those two numbers could rebuild it.

1.2 From a drawing of walls to a graph

A maze on paper is a picture of walls. What an algorithm can generate, store and search is a graph.

That translation is the first decision of this report, and every later section rests on it.

Definition 1 (Area, cells, grid graph). Let $W, H \in \mathbb{N}^*$ be the width and the height of the area, in unit cells. The cells are

$$V = \{(x, y) \in \mathbb{N}^2 : 0 \leq x < W, 0 \leq y < H\}, \quad |V| = WH,$$


and two cells are *adjacent* when they share a side, that is when $|x - x'| + |y - y'| = 1$. The *grid graph* $G_{\text{grid}} = (V, E_{\text{grid}})$ has the cells as vertices and the adjacent pairs as edges, so

$$|E_{\text{grid}}| = \underbrace{H(W - 1)}_{\text{horizontal}} + \underbrace{W(H - 1)}_{\text{vertical}} = 2WH - W - H.$$


Definition 2 (Maze). A *maze* on that area is a spanning subgraph $M = (V, E)$ with $E \subseteq E_{\text{grid}}$. An edge of E is a **door**: the two cells it joins can be walked between. A pair of $E_{\text{grid}} \setminus E$ is a **wall segment**. One distinguished cell v_{exit} is the exit.

Figure 2 shows the two readings of one object.

The vertex set is fixed by W and H alone. So generating a maze means nothing more than **choosing which subset of E_{grid} to keep**.



(a) The maze as walls. The area is enclosed; the exit is a cell, not a gap in the outer wall.



(b) The same maze as a spanning tree of the grid graph.

Figure 2: A maze is not a picture but a graph. Cells become vertices, doors become edges, and a wall is simply an edge of the grid graph that is absent. The exit is a vertex marked on the finished tree — here an interior one, with degree 3 — and neither its position nor its degree affects anything the construction does.

1.3 What the specification forces E to be

Three requirements follow from the statement.

- (R1) **Every cell reaches the exit.** The entry position is arbitrary, so any cell may be a start and each must have a route out. M is connected.
- (R2) **There is exactly one exit.** One cell is designated v_{exit} . It is a *cell*, not a gap in the outer wall: the area is enclosed on every side, and the exit is the one place leading out of it. It may lie anywhere, including strictly inside.
- (R3) **The route out does not fork.** The path to the exit is required to be unique, hence shortest. A cycle in M would give two distinct simple routes between the cells it passes through, so M is acyclic.

Connected and acyclic on a fixed vertex set is a spanning tree. The shape of the answer is settled before any algorithm is written.

Proposition 1 (The object to produce). *A subgraph $M = (V, E)$ of G_{grid} satisfies (R1) and (R3) if and only if it is a spanning tree of G_{grid} . It then has $|E| = WH - 1$ doors, and between any two cells there is exactly one simple path, which is therefore also a shortest one.*

That last clause is what makes Problem 2 well posed. Asking for *the* path out is only meaningful because there is one.

It is also the lever for the rest of the report: Section 2 produces such a tree **by construction**, rather than by checking afterwards.

1.4 What a route costs

A graph alone cannot answer “the shortest way out”. It needs a length.

Walking from a room to an adjacent room through a **door** is **one step**. Every edge of M carries weight $w(e) = 1$.

The length of a route is then the **number of doors it crosses**, and nothing else. Three things follow.

The weights are non-negative.

That is the one precondition Dijkstra’s algorithm asks of a graph [7], so Section 3 applies without further argument.

Distance is depth.

With all weights equal, the cells at distance k from a start form a layer. Section 3 turns on this.

The model has limits, and they are choices.

No weighted doors, so a slow staircase cannot be expressed. No diagonal moves, so the grid is 4-connected — which is what makes the Manhattan distance the natural estimate later, where an 8-connected grid would want the octile distance. No cost for turning.

1.5 Constraints on the exit

Three questions about v_{exit} , none of which constrains the construction.

Chosen before or after generation?

After. Generation takes W and H and nothing else, and no step of it could consult an exit. By Proposition 1 every cell is reachable from every other, so any cell designated afterwards works.

On the boundary?

No. The exit is a cell of the area, not a hole in the wall enclosing it, so nothing obliges it to touch the perimeter.

Exactly one door?

No. Uniqueness of the route out comes from M being a tree, not from the exit’s degree: with 1, 2, 3 or 4 doors there is still exactly one simple path to it from anywhere.

The exit carries no constraint. It is a **label applied to a finished tree**, and every cell is an equally valid candidate. Figure 2 marks an interior one of degree 3.

1.6 The two problems, restated on the graph

Problem 1 — Generation.

Given W , H and a source of random bits, produce a spanning tree M of G_{grid} , drawn so that no two runs need agree, together with the structure recording how it was built. The exit is not an input.

Problem 2 — Escape.

Given M and any start cell $s \in V$, return the unique simple path from s to v_{exit} .

Problem 1 is solved in Section 2, which returns *two* structures rather than one.

Problem 2 is then solved twice, once on each of them: Section 3 searches the graph, Section 4 reads the tree. Section 6 compares what they cost.

1.7 Conventions

Symbol	Meaning
W, H	width and height of the area, in cells
$n = WH$	number of cells, the size of the input
(x, y)	cell in column x (left to right), row y (top to bottom)
$G_{\text{grid}}, E_{\text{grid}}$	the grid graph and its edges: every adjacency
$M = (V, E)$	the maze produced: a spanning tree of G_{grid}
v_{exit}	the single exit cell, anywhere in V
$(x_0, y_0)-(x_1, y_1)$	a region, by its upper-left and lower-right cells

1.8 Scope of the deliverable

The specification fixes the contents, not only the topic: a data structure and what it holds, an algorithm per problem, the size of that structure, and the cost of both algorithms.

Required	Where	State
Data structures, and what they store	Section 2	done
Maze generation algorithm	Section 2	done
Cost of a door, and of a route	Section 1.4	done
Path-finding on the graph	Section 3	done
Path-finding on the tree	Section 4	done
Implementation (optional)	Section 5	done
Size of the structures, cost of every algorithm	Section 6	done

2 Maze generation, and the two structures it returns

Section 1 fixed the target: a spanning tree of G_{grid} , drawn at random.

This section builds one. The structure comes first, because the algorithm is written against it. The guarantee comes last, because the construction rule is what supplies it.

2.1 Two outputs, not one

The generator fills two structures as it goes, and returns both. They are the same maze seen twice, and neither is derived from the other.

The graph: which doors exist.

A set of cells, each holding the neighbours it has a door to. It is what a walker needs, and what Section 3 searches.

The binary tree: how it was built.

One node per division, recording where the wall went and where its door was left, with the two sub-mazes that division created as its children. It is what Section 4 reads.

Building the second costs nothing asymptotically — the recursion already has the information, and only has to keep it — and Section 4 shows it is worth thirty times its price.

2.2 The node record

The maze is a set of nodes, one per cell. A node holds

$$\text{Node} = \left\langle \underbrace{\text{value}}_{\text{its } (x,y) \text{ coordinates}}, \underbrace{\text{neighbours}}_{\text{the cells it has a door to}} \right\rangle.$$

The neighbour list is **empty when the node is created**. A cell has four sides, so it never holds more than four entries, and it grows only when the construction puts a **door** on that node.

An index $\langle P_i, P_j \rangle$ — a $W \times H$ array of references — maps a coordinate pair to its node in $O(1)$. A division reaches the cells it wants without searching. Figure 3 is one node before and after.




Figure 3: The node record. A node carries its coordinates and a list of the neighbours it has a door to; the list is empty when the node is made and is appended to only by a division whose door lands on that node. A cell has four sides, so the list never exceeds four entries, and in the finished maze the entries across all nodes number $2(WH - 1)$ — each door recorded once at each of its two ends.

Each door is recorded at both ends, so the finished maze holds $2(WH - 1)$ entries for $WH - 1$ doors. That buys enumeration of a cell's exits in $O(\deg)$ from the cell alone, which is what Section 3 does at every step.

2.3 Regions, named by two corners

A region is a rectangle of cells, named by the two cells at its ends: (x_0, y_0) – (x_1, y_1) , upper-left and lower-right, with $x_0 \leq x_1$ and $y_0 \leq y_1$. The whole area is $(0, 0)$ – $(W - 1, H - 1)$.

Both corners are **cells of the region**, not fences around it. A region spans $x_1 - x_0 + 1$ columns and $y_1 - y_0 + 1$ rows.

In the pseudocode the two corners are p and q , and a cell's coordinates are reached by index: $p[0]$ is its x , $p[1]$ its y . Indexing rather than naming them lets the division below be written once for both axes.

2.4 The recursion stops at a corridor

Definition 3 (Corridor). A region is a *corridor* when one of its two sides is a single cell: $x_0 = x_1$, one column wide and any height, or $y_0 = y_1$, one row tall and any width. A single cell, where both hold, is the smallest corridor.

A corridor is the case where the construction has nothing to decide.

Its cells lie in a line, so the subgraph G_{grid} induces on it is already a path. A path on m vertices has $m - 1$ edges and is a tree, so it is its own and only spanning tree.

Put a **wall** across a corridor and it has exactly one crossing, which is then forced to be the **door**. The division would change nothing.

So the recursion stops there. **SpanCorridor** lays those forced doors and returns the leaf recording them.




Figure 4: A *corridor* is a region one cell thick, in either direction. Its cells lie in a line, so the induced graph is already a path and there is exactly one way to span it — nothing about a corridor is left to decide, which is why the recursion stops at one. A region thicker than that, on the right, still has both a choice of orientation and a choice of position.

Stopping at a corridor means **every region reaching a division has both sides at least 2**. Both axes always admit a wall, so the orientation is a plain coin with no test in front of it.

2.5 One division contributes exactly one door

Take any door $e = (u, v)$ of the finished maze and walk down the region tree from the root.

At each internal node, u and v either fall in the same child — and the walk continues — or they are separated. They are separated exactly when e crosses that node's wall, and then e is that node's door, being the only crossing there is.

Otherwise the walk ends at a leaf holding both, and e is one of the doors that leaf laid along its corridor.

So the doors split cleanly in two: one per internal node, plus the interiors of the corridor leaves. Nothing is counted twice, nothing is missed.

That licenses the construction to run **forwards**. A division adds its own door and nothing else; a leaf adds its corridor. A wall is never represented at all — it is simply the doors that were never added.


2.6 The second output: the binary tree

The recursion's shape is a binary tree, one node per call.


- An **internal node** records one division: its orientation t , the boundary j the wall sits on, the position i of its door. Its two children are the two regions the division created.
- A **leaf** is a corridor, of one cell or many.

A binary tree with L leaves has $L - 1$ internal nodes. So of the $WH - 1$ doors, $L - 1$ are division doors and $WH - L$ are corridor doors.

Figure 5 shows the first three divisions of a 4×4 area beside the top of the tree they build.



(a) The first three divisions of a 4×4 area.



(b) The top of the region tree it builds. The four boxes at the bottom are regions still to be divided, not corridors.


Figure 5: Every division is one internal node. The node records the orientation t , the wall position j and the door position i it drew, and its two children are the two regions it created; a leaf is a corridor. A tree with L leaves has $L - 1$ internal nodes, so of the $WH - 1$ doors of Proposition 1, $L - 1$ are division doors and the rest are laid down inside the corridors.

2.7 One division step


A call handles a region $(x_0, y_0)-(x_1, y_1)$. If it is a corridor, Section 2.4 disposes of it. Otherwise both sides are at least 2, and:

1. **Toss for the axis t to cut.** Both admit a wall, so the coin is free. Write $t = 0$ for the x axis and $t = 1$ for the y axis; then $1 - t$ is the other one.
2. **Draw the wall position j on axis t ,** uniformly in $[p[t], q[t] - 1]$, **and the door position i on the other axis,** uniformly in $[p[1 - t], q[1 - t]]$.
3. **Link the two cells the door falls on,** appending each to the other's neighbour list.
4. **Recurse on the two regions** — the same corners, cut at j on axis t — and hang them under the node recording $\langle t, j, i \rangle$.

Remark (Departure from the specification). The specification puts the wall across the *longest* dimension. This report tosses for it instead. Correctness is untouched: Proposition 2 never uses which axis was cut, only that one was.



(a) The call links one pair and stops. Every other node is still isolated; no crossing was examined, let alone removed.



(b) After both recursive calls return. Each half is a tree, and the door is still the only edge between them.

Figure 6: One division, with $t = 0$ — the coin named the x axis, so the wall is vertical. The call draws the wall position j on that axis and the door position i on the other, links the two cells the door falls on, and recurses on the two halves. Counting edges: 11 on the left, 7 on the right, one door, total $19 = 20 - 1$ on 20 cells — the tally of Proposition 2.

2.8 The algorithm

Algorithm 1: Maze generation by recursive division.

Input: $W, H \geq 1$, the width and height of the area in cells. The bound is a precondition, not a convention: at $W = 0$ the region $(0, 0) - (-1, H - 1)$ holds no cell and there is nothing for the first draw to land on.

Output: N , the set of nodes holding a spanning tree of G_{grid} ; r , the root of the region tree that built it.

```

1 Function GenerateMaze( $W, H$ ):
2    $N \leftarrow \{ \text{NewNode}(x, y) : 0 \leq x < W, 0 \leq y < H \}$ ;    // every neighbour list empty
3    $r \leftarrow \text{Divide}(N, (0, 0), (W - 1, H - 1))$ ;
4   return  $(N, r)$ ;

5 Function Divide( $N, p, q$ ):
6   if  $p[0] = q[0]$  or  $p[1] = q[1]$  then return SpanCorridor( $N, p, q$ );
7    $t \leftarrow \text{RandomBit}()$ ;                                // the axis to cut: 0 is  $x$ , 1 is  $y$ 
8    $j \leftarrow \text{RandomIn}(p[t], q[t] - 1)$ ;                  // wall between  $j$  and  $j + 1$  on axis  $t$ 
9    $i \leftarrow \text{RandomIn}(p[1-t], q[1-t])$ ;                  // the one door, at  $i$  on the other axis
10   $\text{Link}(N, \text{CellAt}(t, j, i), \text{CellAt}(t, j+1, i))$ ;
11   $\ell \leftarrow \text{Divide}(N, p, \text{WithAxis}(q, t, j))$ ;
12   $r \leftarrow \text{Divide}(N, \text{WithAxis}(p, t, j+1), q)$ ;
13  return CutNode( $t, j, i, \ell, r$ );

14 Function SpanCorridor( $N, p, q$ ):
15   for  $x \leftarrow p[0]$  to  $q[0] - 1$  do  $\text{Link}(N, (x, p[1]), (x+1, p[1]))$ ;
16   for  $y \leftarrow p[1]$  to  $q[1] - 1$  do  $\text{Link}(N, (p[0], y), (p[0], y+1))$ ;
17   return Leaf( $p, q$ );

18 Function Link( $N, u, v$ ):
19    $\text{append } N[v] \text{ to } N[u].\text{neighbours}$ ;
20    $\text{append } N[u] \text{ to } N[v].\text{neighbours}$ ;

21 Function CellAt( $t, u, v$ ):
22   return the cell whose axis- $t$  coordinate is  $u$  and whose other is  $v$ ;

23 Function WithAxis( $p, t, u$ ):
24   return the cell  $p$  with its axis- $t$  coordinate replaced by  $u$ ;
```

`SpanCorridor()` is the whole base case. It lays the corridor's forced doors *and* returns the leaf recording the region.

It needs no case analysis: on a corridor one of its loops is empty by construction, on a single cell both are. The same two loops cover a row, a column and one cell.

Every `Link()` joins two cells of the region that called it, and they are adjacent. So every appended entry is a genuine edge of G_{grid} .

The exit plays no part, for the reasons in Section 1.5. The algorithm never reads it, and any cell may be designated once the tree exists.

2.9 Why the result is a spanning tree

Proposition 2 (Correctness of Algorithm 1). *On a region of m cells, Divide leaves the subgraph induced on that region connected and acyclic, with exactly $m - 1$ doors. In particular GenerateMaze returns a spanning tree of G_{grid} .*

Proof. By induction on m .

Base. If the region has a side of one it is a corridor of m cells in a line, and **SpanCorridor** joins each to the next: a path, connected and acyclic, carrying $m - 1$ doors. This covers $m = 1$, where both loops are empty and $m - 1 = 0$.

Step. Otherwise both sides are at least 2, so $q[0] - p[0] \geq 1$ and $q[1] - p[1] \geq 1$. Whichever axis the coin names, $q[t] - p[t] \geq 1$, so j is drawn from a non-empty range and the two sub-regions R_ℓ and R_r are non-empty, with $m_\ell + m_r = m$.

The call adds the single door at i , whose two cells lie one in each sub-region, then recurses. By the induction hypothesis the two calls leave their regions connected and acyclic with $m_\ell - 1$ and $m_r - 1$ doors. Neither can add an edge between the sub-regions, since every **Link** a call makes joins two cells of its own region.

So the door at i is the only edge between R_ℓ and R_r . Joining two disjoint trees by a single edge yields a tree. It is connected, since any two cells are joined inside their own half or by crossing that door. It is acyclic, since a cycle would have to cross the wall an even, hence non-zero, number of times while only one crossing exists. The door count is $(m_\ell - 1) + (m_r - 1) + 1 = m - 1$.

Applied to the whole area this gives a connected acyclic spanning subgraph of G_{grid} on all WH cells: a spanning tree, with $WH - 1$ doors. $\square$

The guarantee is structural. Nothing is checked after the fact and no attempt is ever rejected and retried.

The uniqueness of the route out — the whole of Problem 2 — follows from the shape of the recursion, not from a property that had to be tested.

2.10 What generation returns

Generation returns both structures: the set N of nodes, each carrying its coordinates and the doors it ended up with, and the root r of the binary tree.

Output	What it holds	Solved by
Graph	each node's neighbours and coordinates	Dijkstra, Section 3
Binary tree	one node per division, with its door	common ancestor, Section 4

The coordinates in the first row are what make the heuristics of Section 3 possible: an estimate of the remaining distance needs to know where the cell is.

Figure 2b is the output of one 4×4 run, whose first three divisions are drawn in Figure 5a: 15 doors on 16 cells, one path between any two of them.

Section 1.4 already gave every door a length of one step, so “the route out” is a quantity. The next two sections find it, once on each output.

3 Solving the graph output: Dijkstra

3.1 Dijkstra on the maze

The first output is a graph of cells joined by doors. Every door weighs 1 by Section 1.4 and no weight is negative, so Dijkstra's algorithm [7] applies to it directly.

It keeps a distance $g(v)$ for every cell, a priority queue of cells reached but not yet settled, and repeatedly settles the nearest one, relaxing its doors.

Two details matter for what follows.

It stops when the exit is *settled*, not when it is reached.

The exit is put in the queue like any other cell; the search ends when it comes back out. Stopping on first sight would be a different algorithm with a different count [11].

On unit weights it is breadth-first search.

With all weights equal, the cell of least g is always one of least depth, so cells are settled in non-decreasing depth order — the layers of Section 1.4. Dijkstra generalises breadth-first search to weighted graphs [4, Ch. 22]; here the generalisation buys nothing but a log factor.

3.2 The two distances to the exit

The search has something Dijkstra does not use: every node carries its coordinates, so the *straight-line* remoteness of a cell from the exit is available for free.

Definition 4 (Estimates). For a cell $v = (x, y)$ and the exit $v_{\text{exit}} = (x_e, y_e)$,

$$h_M(v) = |x - x_e| + |y - y_e|, \quad h_E(v) = \sqrt{(x - x_e)^2 + (y - y_e)^2}.$$

Both are **admissible**: neither can exceed the true number of doors to the exit. Any route changes one coordinate by one per step, so it crosses at least h_M doors on an empty grid, and walls only remove edges, which can only lengthen it. Then $h_E \leq h_M$ by the inequality between the Euclidean and the Manhattan norm.

$h_M \geq h_E$ everywhere, with equality only when the exit shares a row or a column with the cell. **Manhattan dominates Euclidean**: it is the exact distance on a 4-connected empty grid, while the Euclidean value is a straight line no walker can take.

3.3 Two ways to use an estimate

The estimate can enter the search at two different strengths, and the report measures both because they behave very differently.

As a tie-break.

Order the queue by (g, h) . The primary key stays the distance already travelled; h decides only between cells the search has reached at *equal* cost. No computed distance changes.

As a priority.

Order the queue by $(g + h, h)$. This is A* [10]: the estimate is added to the cost, so the search leans towards the exit instead of spreading evenly.

A* does **not** trade the shortest route for a quicker answer. With an admissible estimate it returns an optimal route, exactly as Dijkstra does; it simply settles fewer cells on the way [10]. Indeed Dijkstra *is* A* with $h \equiv 0$ — one algorithm, told nothing.

The methods that do trade optimality for speed are different ones, and are not used here. **Greedy best-first search** orders by h alone, discarding g : fast, and neither optimal nor reliably complete. **Weighted A*** orders by $g + w h$ with $w > 1$, which returns a route at most w times longer than the shortest. Both are worth naming because A* is often remembered as one of them.

3.4 What tie-breaking can and cannot do

The tie-break is the weaker of the two, and by how much can be settled before measuring anything.

Proposition 3 (Ceiling on the tie-break). *Let $d = g(v_{\text{exit}})$. Under any tie-break, the search settles every cell with $g < d$, and then some subset of the cells with $g = d$. Ordering ties by h settles none of that last layer before the exit.*

Proof. The primary key is g , so cells leave the queue in non-decreasing g . Every cell with $g < d$ therefore leaves before v_{exit} , whatever the tie-break: the set $\{v : g(v) < d\}$ is expanded in all cases. The only freedom is the order inside the layer $g = d$. Both estimates are non-negative and vanish exactly at the exit, so $h(v_{\text{exit}}) = 0$ is strictly least in that layer, and the exit is taken first. $\square$

A tie-break cannot prune anything. Its entire effect is bounded by one layer, and the cells with $g < d$ are **necessarily expanded** by any correct search of this shape [6]. It follows too that h_M and h_E must behave **identically** here: both rank the exit first and the search ends there.

Ordering ties by h is the standard $[f, h]$ strategy of the heuristic-search literature [5]; what is unusual here is only that the primary key is blind, so the ceiling above is a corollary of the necessarily-expanded result rather than a new claim.

Using the estimate as a *priority* escapes that ceiling, because it changes f and so changes which cells fall below the threshold at all. A cell deep in a dead end is dropped as soon as its depth plus its estimate exceeds d .

3.5 The algorithm

Algorithm 2: Shortest route to the exit, with an estimate as guide.

Input: the maze M , a start cell s , the exit v_{exit} , an estimate h (possibly $h \equiv 0$), and whether it is used as a tie-break or as a priority.

Output: the route from s to v_{exit} , and the number of cells settled.

```

1 Function ShortestPathToExit( $M, s, v_{\text{exit}}, h, \text{use}$ ):
2    $g[v] \leftarrow \infty$  for every cell;  $g[s] \leftarrow 0$ ;
3    $\text{settled}[v] \leftarrow \text{FALSE}$  for every cell;
4    $Q \leftarrow$  empty priority queue; Push( $Q, s$ );
5    $\text{expanded} \leftarrow 0$ ;
6   while  $Q$  is not empty do
7      $u \leftarrow \text{PopLeast}(Q)$ ;                                // least  $(g + h, h)$ , or least  $(g, h)$ 
8     if  $\text{settled}[u]$  then continue;
9      $\text{settled}[u] \leftarrow \text{TRUE}$ ;  $\text{expanded} \leftarrow \text{expanded} + 1$ ;
10    if  $u = v_{\text{exit}}$  then break;
11    for  $v \in M.\text{neighbours}(u)$  do
12      if  $g[u] + 1 < g[v]$  then
13         $g[v] \leftarrow g[u] + 1$ ;                                // one door, one step
14         $\text{parent}[v] \leftarrow u$ ; Push( $Q, v$ );
15  return RouteBackFrom( $\text{parent}, v_{\text{exit}}$ ),  $\text{expanded}$ ;
```

The queue is ordered by $g + h$ when the estimate is a priority and by g when it is a tie-break, with h separating equals in both cases. Setting $h \equiv 0$ recovers plain Dijkstra, so one procedure covers every row of the table in Section 5.

expanded is the figure that section reports. Every setting returns the same route — Proposition 1 leaves only one — so it is the only thing there is to compare.

It is also what Section 4 improves on, by not searching the graph at all.

4 Solving the tree output: the common ancestor

Dijkstra settles most of the maze. The second output makes that unnecessary.

4.1 Addresses

Every cell lies in exactly one leaf of the binary tree, and the sequence of turns from the root to that leaf is the cell's **address** (Figure 7). Deciding which side of a division a cell falls on is a comparison of its coordinates against a rectangle — no cell of the maze is read, no neighbour list consulted.




Figure 7: The proposed structure. Each internal node stores the branching point of the division that created it — the cell its single door sits on — and each leaf is a region of the maze. A cell's **address** is the sequence of turns from the root to its leaf, so RRL names one region in three characters. Sibling leaves are the two halves a single division separated, and are joined in the maze by that division's door alone.

4.2 The lemma that does the work



Proposition 4 (Separation). *Let R be an internal node, R_ℓ and R_r its children, and (u, v) its door with $u \in R_\ell$ and $v \in R_r$. For cells a, b of R , the unique route from a to b crosses (u, v) if and only if a and b lie in different children.*

Proof. Proposition 2 established that (u, v) is the only edge between R_ℓ and R_r , and that each child is spanned by a tree of its own.

If a and b lie in different children, a route between them uses an edge between the children, and (u, v) is the only one.

If both lie in the same child, that child's own spanning tree already contains a route between them. The maze is a tree, so its route between two cells is unique and is therefore that one; it stays inside the child. $\square$

The right-hand case is where the saving lives. When both endpoints fall on one side, **half the region is discarded in a single test** — and the test never looks at the maze.



Size	Reached the exit	Never terminated	Stuck
16×16	13/200	187	—
32×32	6/200	194	—
64×64	2/200	198	—

Instrumenting the 32×32 failures finds 199 stopped at a local minimum and 0 cycled (Figure 9). Two reasons.

It is greedy best-first search.

Ranking moves by an estimate of the remaining distance while ignoring the distance travelled is neither optimal nor, without a closed list, complete.

A node’s door is not where the walker is.

An internal node stands for a *region*; its parent covers the same ground and more. Ranking a parent against a child by one stored cell compares things that are not comparable.




Figure 9: A local minimum, which is where the greedy rule ends 199 runs in 200. The stored branching point of the current node happens to lie nearer the exit than those of its parent, its two children and its sibling; no move improves h , and the exit is elsewhere. Nothing about the maze is wrong — the rule is simply hill-climbing on a landscape with hills in it.

The estimate answers “how far is this node’s door from the exit”. Algorithm 3 asks “which side holds the exit” — a question the tree answers exactly, with no arithmetic at all.

5 Implementation in C

The subject leaves code optional and provides a skeleton in C. It is written here because several claims in this report are the kind that should be measured rather than asserted: that generation always returns a spanning tree, that both outputs return the same route, and that reading the tree is cheaper than searching the graph.

5.1 What was built

C11, one Makefile, no dependency outside the standard library, compiled with `-Wall -Wextra -Wpedantic -Werror`. It lives beside this report in `implementation/`.

File	Role
<code>cell.{h,c}</code>	a cell, and the two axis operations of Section 2.3
<code>maze.{h,c}</code>	the first output: nodes, coordinates, neighbours
<code>region_tree.{h,c}</code>	the second output: divisions, doors, addresses
<code>generator.{h,c}</code>	Algorithm 1, filling both
<code>solver.{h,c}</code>	Algorithm 2 on the graph, Algorithm 3 on the tree
<code>invariants.{h,c}</code>	the properties a finished maze must have
<code>svg.{h,c}</code>	drawing a maze, and a route through it
<code>tests.c</code>	the checks below
<code>benchmark.c</code>	the measurements of Section 6

Two decisions are worth naming. The invariants are their own module, because a generator that checked its own work would only ever confirm its own assumptions. And the random numbers come from a small `xorshift` generator rather than `rand()`, so that a seed determines a maze on any machine and the measurements are reproducible.

5.2 What is checked

1121 assertions, all passing.

Generation returns a spanning tree.

$WH - 1$ doors, connected, no cell with more than four, over ten shapes $\times$ twenty seeds including 1×1 , 1×40 and 40×1 .

A corridor is laid down whole.

A 1×40 area gives a tree of one leaf and 39 doors — Section 2.4 observed rather than argued.

The same seed gives the same maze.

Without which no measurement below is reproducible.

Both outputs agree, cell for cell.

The graph search and the tree search return identical routes on every shape and seed tried. This is the claim the whole of Section 4 rests on.

Every guide returns the same route.

Plain Dijkstra and all four heuristic settings agree, which is Proposition 1 made visible.

The route is walkable.

Consecutive cells on it are joined by a door.

Bad input is refused.

A zero dimension, a door between cells that do not touch, a door opened twice, a cell outside the grid.

The fourth item is the one that matters. Two algorithms that share no code — one walking the graph with a priority queue, one descending the tree with rectangle tests — return **the same cells in the same order**, every time.

6 Time analysis and comparison

Write $n = WH$ for the number of cells.

6.1 Size of the two structures

Output	What	How many
Graph	nodes	n
	neighbour entries	$2(n - 1)$
	index $\langle P_i, P_j \rangle$	n
Tree	nodes	$2L - 1$, for L corridor leaves

Both are $\Theta(n)$, and the tree is the smaller of the two: L measured between 54 and 8843 on areas of 256 to 65 536 cells, so well under n . Keeping the second output costs a constant factor of memory and no asymptotics.

6.2 Cost of generation

Each call of `Divide` does $O(1)$ work of its own; each call of `SpanCorridor` does work proportional to its corridor. Every door is opened exactly once, and there are $n - 1$ of them.

Generation runs in $\Theta(n)$ time and fills *both* outputs while doing it. This is optimal: any generator must write $n - 1$ doors.

The recursion is not balanced — the wall position is uniform, not central, the same skew that makes quicksort’s worst case quadratic. Measured depth against a balanced tree on the same leaves:

Area	Corridor leaves	Mean depth	Depth if balanced
16×16	54	11.5	5.8
64×64	666	21.4	9.4
256×256	8843	32.0	13.1

About 2.4 times a balanced tree at every size — the figure a stack-depth bound has to be written against.

6.3 Cost of solving the graph

The maze is a tree, so $|E| = n - 1$. With a binary heap Dijkstra costs $O((|V| + |E|) \log |V|) = O(n \log n)$ [4, Ch. 22], and the log buys nothing: all weights are 1, so it settles cells in depth order and is breadth-first search wearing a priority queue.

6.4 Cost of solving the tree

At each division either both endpoints stay on one side, and half the region is discarded, or they are separated and both halves are entered. Write p for the probability of the second. Each call doing $O(1)$ work,

$$T(m) = (1 - p)T(m/2) + p \cdot 2T(m/2) + O(1) = (1 + p)T(m/2) + O(1),$$

so by the master theorem

$$T(n) = \Theta\left(n^{\log_2(1+p)}\right).$$

The recurrence has one parameter, so it makes one prediction. Measuring p directly, 200 mazes per size:

n	One-way	Two-way	p	$\log_2(1+p)$	Exponent fitted to the visits
256	10	19	0.65	0.72	0.70
1024	27	44	0.62	0.69	0.69
4096	69	111	0.62	0.69	0.68
16384	170	266	0.61	0.69	0.68
65536	435	670	0.61	0.68	0.67

Prediction and measurement agree at every size. The two ends of the scale are worth reading off: at $p = 1$ the recurrence is $2T(n/2)$ and the cost is $\Theta(n)$, no better than Dijkstra; at $p = 0$ it is $T(n/2)$ and the cost is $\Theta(\log n)$.

$\Theta(\log n)$ is not available. It would need the endpoints never to be separated, which happens only when they share a corridor. $\log n$ is the cost of **one** root-to-leaf descent, and the search pays one per separation.

6.5 The comparison

200 mazes per size, start at $(0, 0)$, exit at the far corner, counting cells settled by the graph search and nodes visited by the tree search.

Size	Cells	Steps	graph output: cells examined				tree
			Dijkstra	tie: M	A*: M	A*: E	nodes
16×16	256	64	222	221	175	191	49
32×32	1024	180	873	870	738	776	116
64×64	4096	538	3533	3530	3199	3280	293
128×128	16384	1564	13939	13934	12983	13199	702
256×256	65536	4493	56193	56186	53753	54276	1776

All 1000 mazes were checked to be spanning trees. All of them were.




Figure 10: Both axes logarithmic, so a power law is a straight line and the exponent is its slope. Dijkstra runs parallel to n ; the region-tree search runs at slope 0.68, below even the length of the route it returns, because one visit to a corridor leaf covers a whole straight run.

Three conclusions, in order of size.

The tree beats the graph by a factor of thirty.

1776 against 56 193 at 256×256 , and the gap widens with n because the exponents differ:

$\Theta(n^{0.68})$ against $\Theta(n)$. It also wins on memory — $O(\log n)$ of recursion stack against $\Theta(n)$ of arrays Dijkstra must allocate before it starts.

Heuristic tie-breaking does nothing.

5 cells saved in 56 193, and Manhattan and Euclidean are *identical* in every row. Both are predicted: cells nearer the start than the exit are necessarily expanded [6], so only the last layer is negotiable, and both estimates put the exit first in it.

A* helps a little, and Manhattan beats Euclidean.

53 753 against 56 193, a 4% saving, with Manhattan ahead at every size as its dominance predicts. Still 82% of the maze.

6.6 Why the heuristic is weak here, and the tree is not

The two results above have one explanation between them.

A* prunes *rival routes* — on an open grid, the many equal-length ways to reach the same cell. A perfect maze is a tree and has none, so all that is left for A* to prune is dead-end branches whose depth plus estimate exceeds the answer. And the estimate itself is poor: Sturtevant et al. [14] measure that on 512×512 maze maps the geometric estimate averages 151 where the true distance is 512–768, and that a memory-based heuristic expands over eleven times fewer nodes.

A straight line says little about a route that must wind. The tree search wins because it does not estimate anything: it **knows** where the walls are, having built them.

For completeness: the grid-pathfinding methods with the largest speedups, Jump Point Search and the symmetry-breaking family [9], work by collapsing equal-cost alternative routes. A perfect maze has no path symmetry at all, so they have nothing to prune here. Standard benchmark sets include maze maps for that reason [13].

6.7 Where recursive division sits

The maze-generation literature is small and largely practitioner-written; the distinction is marked because it bears on how much weight each claim carries.

The vocabulary. What this report builds is a **perfect maze** — any two cells joined by a unique path — which is the same object as a spanning tree of the grid graph. The term is Pullen’s [12] and is used verbatim in peer-reviewed work [2]. Recursive division is a **wall adder**: it places walls where most generators carve passages.

The alternatives. Randomised depth-first search, randomised Kruskal and Prim, Eller’s algorithm, Hunt-and-Kill, binary tree and sidewinder all produce spanning trees of the same grid in $\Theta(n)$ or close to it. Two do something this one does not: Aldous–Broder [1, 3] and Wilson’s algorithm [15] sample a **uniform** spanning tree, every perfect maze equally likely. Wilson’s runs in expected mean hitting time rather than cover time, a $\Theta(\log n)$ improvement on a grid.

Recursive division is not uniform. Its first division leaves exactly one door crossing a full straight line, and recursively so; a spanning tree admitting no such decomposition can never be produced. The support is a strict subset, so the distribution is not uniform before one even asks how it is weighted inside it. Pullen [12] classifies it accordingly and records its visible signature: long straight walls crossing the interior, rectangular rooms, a self-similar texture. An n -cell grid has on the order of 3.21^n spanning trees, and recursive division reaches only those with a guillotine decomposition.

Recursive division was chosen for being **simple, linear, recursive** — **and for the second output it gives free**. It was not chosen for the quality of its randomness, and should not be claimed for it.

What is missing from the literature. No peer-reviewed analysis of the distribution recursive division induces appears to exist. Bellot et al. [2] include it but score only difficulty; Gabrovšek [8] measures six generators and omits it. Measuring its degree distribution against the uniform spanning tree would close that gap, and is the natural continuation of this work.

References

- [1] David J. Aldous. The random walk construction of uniform spanning trees and uniform labelled trees. *SIAM Journal on Discrete Mathematics*, 3(4):450–465, 1990.
- [2] Valérie Bellot, Maxime Cautrès, Jean-Marie Favreau, Manon Gonzalez-Thauvin, Pascal Lafourcade, Kergann Le Cornec, Maëlig Mosnier, and Sol Rivière-Wekstein. How to generate perfect mazes? *Information Sciences*, 572:444–459, 2021.
- [3] Andrei Z. Broder. Generating random spanning trees. In *30th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 442–447, 1989.
- [4] Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms*. MIT Press, 4 edition, 2022.
- [5] Augusto B. Corrêa, André G. Pereira, and Marcus Ritt. Analyzing tie-breaking strategies for the A* algorithm. In *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence (IJCAI)*, pages 4715–4721, 2018.
- [6] Rina Dechter and Judea Pearl. Generalized best-first search strategies and the optimality of A*. *Journal of the ACM*, 32(3):505–536, 1985.
- [7] E. W. Dijkstra. A note on two problems in connexion with graphs. *Numerische Mathematik*, 1:269–271, 1959.
- [8] Peter Gabrovšek. Analysis of maze generating algorithms. *IPSI Transactions on Internet Research*, 15(1):23–30, 2019.
- [9] Daniel Harabor and Alban Grastien. Online graph pruning for pathfinding on grid maps. In *Proceedings of the Twenty-Fifth AAAI Conference on Artificial Intelligence*, volume 25, pages 1114–1119, 2011.
- [10] Peter E. Hart, Nils J. Nilsson, and Bertram Raphael. A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*, SSC-4(2):100–107, 1968.
- [11] Robert C. Holte. Common misconceptions concerning heuristic search. In *Proceedings of the Third Annual Symposium on Combinatorial Search (SoCS)*, pages 46–51, 2010.
- [12] Walter D. Pullen. Think labyrinth: Maze algorithms. <https://www.astrolog.org/labyrinth/algrithm.htm>, 2026. Practitioner reference, not peer-reviewed; the source of the field’s vocabulary and its standard comparison table.
- [13] Nathan R. Sturtevant. Benchmarks for grid-based pathfinding. *IEEE Transactions on Computational Intelligence and AI in Games*, 4(2):144–148, 2012.
- [14] Nathan R. Sturtevant, Ariel Felner, Max Barrer, Jonathan Schaeffer, and Neil Burch. Memory-based heuristics for explicit state spaces. In *Proceedings of the Twenty-First International Joint Conference on Artificial Intelligence (IJCAI)*, pages 609–614, 2009.
- [15] David Bruce Wilson. Generating random spanning trees more quickly than the cover time. In *Proceedings of the Twenty-Eighth Annual ACM Symposium on Theory of Computing (STOC)*, pages 296–303, 1996.